

Data Warehousing

Lecture 10 Density and Partition based Data Clustering

CITS3401
CITS5504

Zeyi Wen

Computer Science and
Software Engineering

School of Maths, Physics
and Computing

Acknowledgement: The lecture slides are based on online sources.

Project 2: Common Questions

- The data have been modified by the data set provider (e.g. to hide some privacy).
 - It is like performing normalisation.
 - We can normalise any attribute to the range of $[0,1]$, where 0 just represents the smallest value of that attribute.
- You can decide how to deal with those “nonsensible” attributes.
 - Remove them, use them as they are, transform them, etc.
- Do I need to discretise the attributes?
 - Better to do that for association rule mining
 - May or may not need discretisation in other tasks.

- Introduction to Clustering
 - Computing Distance: Review of Lecture 7
- Density based Clustering
 - DBScan
- Partition based Clustering
 - K-Means
 - K-Medoids

What is Cluster Analysis?


- **Cluster:**
 - instances similar to one another are within the same cluster
 - Instances dissimilar are in different clusters
- **Cluster analysis:** Finding characteristics for similar instances
- **Unsupervised learning:** no predefined classes
- **Typical applications**
 - As a **stand-alone tool** to get insight into data distribution
 - As a **preprocessing step** for other algorithm
- **Rich Applications**
 - Document classification
 - Market research
 - DNA analysis
 - Create thematic maps in GIS

What is good clustering?

- Good clustering methods produce high quality clusters with
 - high intra-class similarity
 - low inter-class similarity
 - The definitions of similarity, measured as a **distance functions** may be different for numeric, boolean, and categorical attributes. Often is highly problem dependent.
- The quality of a clustering method is also measured by its **ability to discover some or all of the hidden patterns.**

Major Types of Clustering Algorithms

- **Partition-based Clustering**
 - K-Means Algorithm
 - K-Medoids Algorithm
- **Density-Based Clustering**
 - DBScan



Major Clustering Approaches

- **Partitioning approach:** *k-means, k-medoids, CLARANS*
 - Construct k -partitions for the given n -instances ($k \leq n$). Each group contains at least one instance. Each instance must belong to exactly one group.
- **Hierarchical approach:** *Diana, Agnes, BIRCH, ROCK, CAMELEON*
 - Create a hierarchical decomposition of the set of objects using some criterion (linkage function)
 - Agglomerative Approach: bottom-up merging
 - Divisive Approach: top-down splitting
- **Density-based approach:** *DBSACN, OPTICS, DenClue*
 - Based on connectivity and density functions. i.e. for each data point within a given cluster, the radius of a given cluster has to contain at least a minimum number of points.

- Introduction to Clustering
 - Computing Distance: Review of Lecture 7
- Density based Clustering
 - DBScan
- Partition based Clustering
 - K-Means
 - K-Medoids

- Distances are normally used to measure the **similarity** or **dissimilarity** between two data objects
- Some popular ones include: Minkowski distance:

- $$d(i, j) = \sqrt[q]{|x_{i1} - x_{j1}|^q + |x_{i2} - x_{j2}|^q + \cdots + |x_{ip} - x_{jp}|^q}$$

- where $i = (x_{i1}, x_{i2}, \dots, x_{ip})$ and $j = (x_{j1}, x_{j2}, \dots, x_{jp})$ are two p -dimensional instances, and q is a positive integer.

- If $q=1$, d is **Manhattan** Distance

- $$d(i, j) = |x_{i1} - x_{j1}| + |x_{i2} - x_{j2}| + \cdots + |x_{ip} - x_{jp}|$$

- If $q=2$, d is **Euclidean** Distance

- $$d(i, j) = \sqrt{|x_{i1} - x_{j1}|^2 + |x_{i2} - x_{j2}|^2 + \cdots + |x_{ip} - x_{jp}|^2}$$

- Properties:

- $d(i, j) \geq 0$; $d(i, j) = 0$; $d(i, j) = d(j, i)$ and

- $d(i, j) \leq d(i, k) + d(j, k)$; (**Triangle inequality**)

- Also, one can use weighted distance, parametric Pearson product moment correlation, or other dissimilarity measures

Compute Dissimilarity for Binary Attributes

- **A contingency table for binary data**
 - Symmetric: both matches are equally important.
 - Asymmetric: the match of “0” is not important.

- **Distance Measure for**

- Symmetric binary attributes:

- $$d(i, j) = \frac{b+c}{a+b+c+d} = \frac{b+c}{p}$$

- Asymmetric binary attributes:

- $$d(i, j) = \frac{b+c}{a+b+c}$$

		Instance j		
Instance i		1	0	total
	1	a	b	a + b
	0	c	d	c + d
	total	a + c	b + d	p

- **Jaccard Coefficient (similarity measure for asymmetric binary attributes):**

- $$sim_{Jaccard}(i, j) = \frac{a}{a+b+c}$$

- **Convert binary attribute into numerical attribute**

Dissimilarity between Asymmetric Binary Attributes

- Given the following example

Name	Gender	Fever	Cough	Test-1	Test-2	Test-3	Test-4
Jack	M	Y	N	P	N	N	N
Mary	F	Y	N	P	N	P	N
Jim	M	Y	P	N	N	N	N

	Mary		
	1	0	Σ_{row}
Jack	1	0	2
	0	3	4
Σ_{col}	3	3	6

- Gender is a **symmetric** attribute (**not counted in**)
- The remaining attributes are **asymmetric** binary
- Let the values of Y and P to be 1, and value N to be 0.

- We have

$$d(jack, mary) = \frac{0+1}{2+0+1} = 0.33$$

$$d(jack, jim) = \frac{1+1}{1+1+1} = 0.67$$

$$d(jim, mary) = \frac{1+2}{1+1+2} = 0.75$$

	Jim		
	1	0	Σ_{row}
Jack	1	1	2
	0	3	4
Σ_{col}	2	4	6

	Mary		
	1	0	Σ_{row}
Jim	1	1	2
	0	2	4
Σ_{col}	3	3	6

- A generalisation of the binary variable in that it can take more than 2 states, e.g. red, yellow, blue, green
- Method 1: Simple matching
 - m : # of matches, p : total # of attributes
 - $d(i, j) = \frac{p-m}{p}$
- Method 2: convert to a number of binary attributes
 - creating a new binary attribute for each of the M possible states

- An ordinal attribute is often discrete.
- Order is important, e.g. rank (e.g. freshman, sophomore)
- Can be treated as numeric attributes
 - replace x_{if} by their rank $r_{if} \in \{1, \dots, M_f\}$
 - map the range of each attribute onto $[0, 1]$ by replacing i -th object in the f -th attribute by
 - $z_{if} = (r_{if} - 1)/(M_{if} - 1)$
 - example: freshman: 0; sophomore: 1/3; junior: 2/3; senior 1
 - distance: $d(\text{freshman}, \text{senior}) = 1$, $d(\text{junior}, \text{senior}) = 1/3$
 - compute the dissimilarity using methods for numeric attributes

- A database may contain **different types** of attributes
 - symmetric binary, asymmetric binary, nominal, ordinal, and numeric attributes
- One may use a weighted formula to combine their effects

$$d(i, j) = \frac{\sum_{f=1}^p \delta_{ij}^f d_{ij}^f}{\sum_{f=1}^p \delta_{ij}^f}$$

- f is binary or nominal:

$$\begin{aligned} d_{ij}^f &= 0, & \text{if } x_{if} = x_{jf} & \text{ or} \\ d_{ij}^f &= 1, & \text{otherwise} \end{aligned}$$

- f is numeric: use Euclidean distance
- f is ordinal

- compute ranks z_{if} where $z_{if} = (r_{if} - 1)/(M_{if} - 1)$
- and treat z_{if} as numeric attribute

Similarity of Two Vectors

- Vector objects: keywords in documents, gene features in micro-arrays, etc.

<i>Document</i>	<i>team</i>	<i>coach</i>	<i>hockey</i>	<i>baseball</i>	<i>soccer</i>	<i>penalty</i>	<i>score</i>	<i>win</i>	<i>loss</i>	<i>season</i>
Document1	5	0	3	0	2	0	0	2	0	0
Document2	3	0	2	0	1	1	0	1	0	1
Document3	0	7	0	2	1	0	0	3	0	0
Document4	0	1	0	0	1	2	2	0	3	0

- Broad applications: information retrieval, natural language understanding, etc.
- Cosine measure

$$- s(x, y) = \frac{d_1^T \cdot d_2}{\|d_1\| * \|d_2\|}$$


- A variant: Tanimoto coefficient-used in information retrieval

$$- s(x, y) = \frac{d_1^T \cdot d_2}{d_1^T \cdot d_1 + d_2^T \cdot d_2 - d_1^T \cdot d_2}$$

- Introduction to Clustering
 - Computing Distance: Review of Lecture 7
- **Density based Clustering**
 - DBScan
- Partition based Clustering
 - K-Means
 - K-Medoids

- **Clustering based on density (local cluster criterion), such as density-connected points**
- **Major features:**
 - Discover clusters of arbitrary shape
 - Handle noise
 - **One scan**
 - Need density parameters as termination condition
- **Several interesting studies:**
 - DBSCAN: Ester, et al. (KDD'96)
 - OPTICS: Ankerst, et al (SIGMOD'99).
 - DENCLUE: Hinneburg & D. Keim (KDD'98)
 - CLIQUE: Agrawal, et al. (SIGMOD'98) (more grid-based)

- Two parameters:
 - *Eps*: Maximum radius of the neighbourhood
 - *MinPts*: Minimum number of points in an Eps-neighbourhood of that point
- $N_{Eps}(p): \{q \in D \mid dist(p, q) \leq Eps\}$
- **Directly density-reachable**: A point p is directly density-reachable from a point q w.r.t. Eps , $MinPts$ if
 - p belongs to $N_{Eps}(q)$ & q is a core.
 - Core point condition:
 - $|N_{Eps}(q)| \geq MinPts$




MinPts = 5

Eps = 1 cm



- **Density-reachable:**

- A point p is **density-reachable** from a point q w.r.t. Eps , $MinPts$ if there is a chain of points $p_1, \dots, p_n, p_1 = q, p_n = p$ such that p_{i+1} is directly density-reachable from p_i




- **Density-connected**

- A point p is **density-connected** to a point q w.r.t. Eps , $MinPts$ if there is a point o such that both, p and q are density-reachable from o w.r.t. Eps and $MinPts$



DBSCAN: Density Based Spatial Clustering of Applications with Noise

- Relies on a *density-based* notion of cluster: A *cluster* is defined as a maximal set of density-connected points
- Discovers clusters of arbitrary shape in spatial databases with noise



DBSCAN: The Algorithm

- Arbitrary select a point p
- Retrieve all points density-reachable from p w.r.t. Eps and $MinPts$.
- If p is a core point, a cluster is formed.
- If p is a border point, no points are density-reachable from p and DBSCAN visits the next point of the database.
- Continue the process until all of the points have been processed.

DBSCAN: Sensitive to Parameters

Figure 8. DBScan results for DS1 with MinPts at 4 and Eps at (a) 0.5 and (b) 0.4.

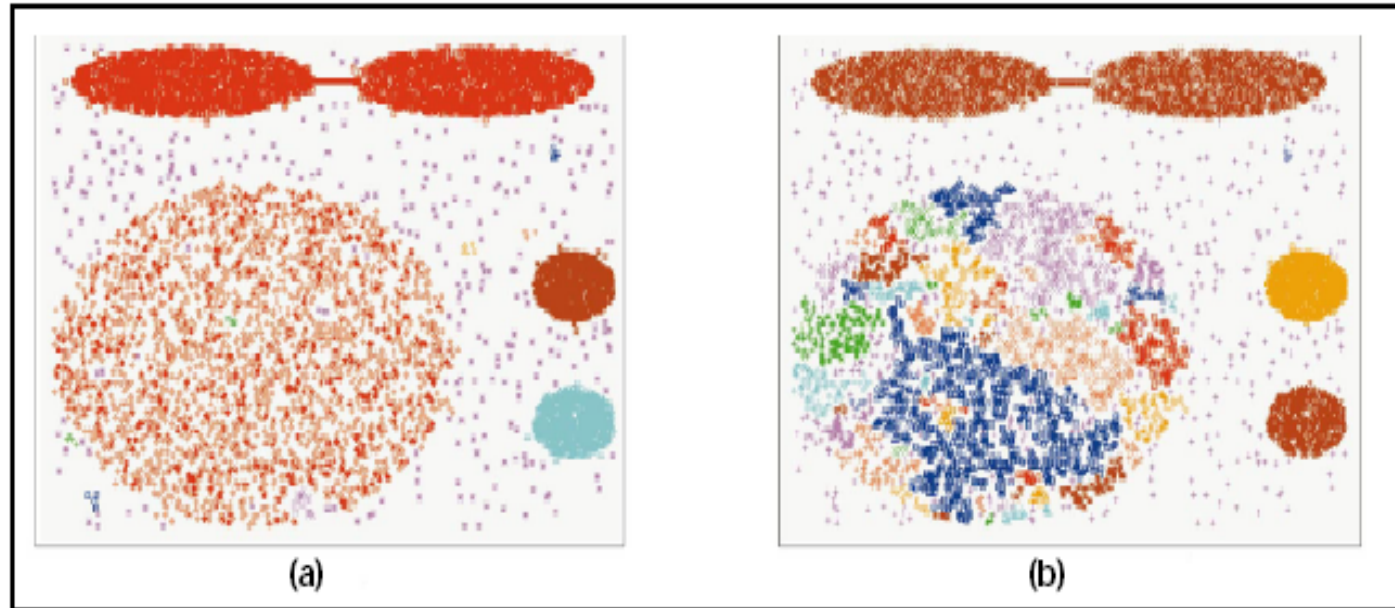



Figure 9. DBScan results for DS2 with MinPts at 4 and Eps at (a) 5.0, (b) 3.5, and (c) 3.0.




- Introduction to Clustering
 - Computing Distance: Review of Lecture 7
- Density based Clustering
 - DBScan
- **Partition based Clustering**
 - K-Means
 - K-Medoids

- **Partitioning method:** Construct a partition of a database D of n instances into a set of k clusters, s.t., minimise the sum of squared distance (within cluster variance)
 - $E = \sum_{i=1}^k \sum_{p \in C_i} (p - m_i)^2$
- Given an integer k , find a partition of k clusters that optimises the chosen partitioning criterion
 - Global optimal: **exhaustively enumerate all partitions**
 - Heuristic methods: *k-means* and *k-medoids* algorithms
 - **k-means** (MacQueen'67): Each cluster is represented by the **center** of the cluster
 - **k-medoids** or PAM (Partition around medoids) (Kaufman & Rousseeuw'87): Each cluster is represented by **one of the instances** in the cluster

Given k , the *k*-means algorithm is implemented in four steps:


1. Partition instances into k nonempty subsets
2. Compute seed points as the centroids of the clusters of the current partition (the centroid is the center, i.e. *mean point*, of the cluster)
3. Assign each object to the cluster with the nearest seed point
4. Go back to Step 2, stop when no more new assignment

Example: the k-Means Clustering Method




$K=2$
Arbitrarily choose K
instances as initial
cluster center


Assign
each
instances
to most
similar
center



reassign




Update
the
cluster
means



reassign

Update
the
cluster
means



K-means clustering algorithm

- **Input:** K , a set of points x_1, x_2, \dots, x_n where $x_i \in \mathbb{R}^d$
- Initialise K centroids c_1, c_2, \dots, c_K at random locations
- Repeat until convergence

→ for each point x_i : **Assignment step**

* find the nearest centroid c_j

* assign x_i to cluster j


$$c_{j,a} = \frac{1}{n_j} \sum_{x_i \rightarrow c_j} x_{i,a} \quad \text{for } a \in \{1, 2, \dots, d\}$$

→ for each centroid c_j : **Update step**


* update c_j to the mean of all the points assigned to cluster j

- “Convergence” means none of the points changes its cluster.


K-means clustering example




K-means clustering example




K-means clustering example




K-means clustering example




K-means clustering example




K-means clustering example




K-means clustering example




K-means clustering example




K-means clustering example




K-means clustering example



- Strength: Relatively efficient: $O(tkn)$, where n is # instances, k is # clusters, and t is # iterations.
 - Normally, $k, t \ll n$.
- Often terminates at a **local optimum**. The global optimum may be found using techniques such as:
 - deterministic annealing and genetic algorithms
- **Weakness**
 - Applicable only **when mean is defined**, then what about categorical data?
 - **Need to specify k** , the number of clusters, in advance
 - Unable to handle **noisy data and outliers**
 - Not suitable to discover clusters with **non-convex shapes**



Convex Polygons



Non-convex Polygons

- A few variants of the *k-means* which differ in
 - Choosing better initial *centroids*
 - *k-means++*, *Intelligent k-means*, *Genetic k-means*
 - Dissimilarity calculations
 - Strategies to calculate cluster means
- Handling categorical data: *k-modes* (Huang'98, **aside**)
 - Using new dissimilarity measures to deal with categorical attributes
 - A mixture of categorical and numerical data: *k-prototype* method
 - Replacing means of clusters with modes
 - Using a frequency-based method to update modes of clusters

Centroid, Radius and Diameter of a Cluster (for numerical data sets)

- Centroid: the “middle” of a cluster
 - $C_m = \frac{\sum_{i=1}^N t_{ip}}{N}$
- Radius: square root of average distance from any point of the cluster to its centroid
 - $R_m = \sqrt{\frac{\sum_{i=1}^N (t_{ip} - C_m)^2}{N}}$
- Diameter: square root of average mean squared distance between all pairs of points in the cluster

- $D_m = \sqrt{\frac{\sum_{i=1}^N \sum_{j=1}^N (t_{ip} - t_{jq})^2}{N(N-1)}}$



Calculate the distance between Clusters

- **Single link**
 - smallest distance between an element in one cluster and an element in the other, i.e. $dis(K_i, K_j) = \min(t_{ip}, t_{jp})$
- **Complete link**
 - largest distance between an element in one cluster and an element in the other, i.e. $dis(K_i, K_j) = \max(t_{ip}, t_{jp})$
- **Average**
 - average distance between an element in one cluster and an element in the other, i.e. $dis(K_i, K_j) = avg(t_{ip}, t_{jp})$
- **Centroid**
 - distance between the centroids of two clusters, i.e. $dis(K_i, K_j) = dis(C_i, C_j)$
- **Medoid:**
 - medoid is the most centrally located object in a cluster
 - distance between the medoids of two clusters, i.e. $dis(K_i, K_j) = dis(M_i, M_j)$

- Introduction to Clustering
 - Computing Distance: Review of Lecture 7
- Density based Clustering
 - DBScan
- Partition based Clustering
 - K-Means
 - **K-Medoids**

The problem of k-means method


- **The k-means algorithm is sensitive to outliers!**
 - Since an instance with an extremely large value may substantially distort the distribution of the data.
- **K-Medoids:**
 - Instead of taking the mean value of the instance in a cluster as a reference point, medoids can be used, which is the most centrally located instance in a cluster.




The *K-Medoids* Clustering Method

- Find *representative* instances, called medoids, in clusters
- ***PAM* (Partitioning Around Medoids, 1987)**
 - starts from an initial set of medoids and **iteratively replaces** one of the medoids by one of the non-medoids if it **improves the total distance** of the resulting clustering
 - *PAM* works effectively for small data sets, but does not scale well for large data sets
- ***CLARA* (Kaufmann & Rousseeuw, 1990)**
- ***CLARANS* (Ng & Han, 1994): Randomised sampling**
- **Focusing + spatial data structure (Ester et al., 1995)**


A Typical K-Medoids Algorithm (PAM)



Arbitrary
choose k
object as
initial
medoids



Assign
each
remainin
g object
to
nearest
medoids




$K=2$


**Do loop
Until no
change**

Swapping O
and O_{random}
If quality is
improved.

Total Cost = 26




Compute
total cost of
swapping




- **Use a real instance/object to represent the cluster**
 1. Select k representative instances arbitrarily
 2. For each pair of non-selected instance h and selected instance i , calculate the total swapping cost TC_{ih}
 3. For each pair of i and h ,
 - If $TC_{ih} < 0$, i is replaced by h
 - Then assign each non-selected instance to the most similar representative instance
 4. repeat steps 2-3 until there is no change

Determining O_{random} a good replacement of O_j


- We calculate the distance from every object p to the closest object in the set $\{O_1, \dots, O_{j-1}, O_{random}, O_{j+1}, \dots, O_k\}$,
- Then use the distance to update the cost function.
 - Suppose object p is currently assigned to a cluster represented by medoid O_j (Figure a or b). Do we need to reassign p to a different cluster (represented by medoid O_i) if O_j is being replaced by O_{random} ?
 - What if object p is currently assigned to a different cluster represented by medoid O_i ? (Figure c or d)




(a) Reassigned
to O_i



(b) Reassigned
to O_{random}



(c) No change




(d) Reassigned
to O_{random}


- Data object
- + Cluster center
- Before swapping
- After swapping

PAM Clustering: Total Swapping Cost: $TC_{ih} = \sum_j C_{jih}$


Replace i by h , the following scenarios are the distance change on j to its nearest medoid




$$C_{jih} = d(j, h) - d(j, i)$$



$$C_{jih} = 0$$



$$C_{jih} = d(j, t) - d(j, i)$$



$$C_{jih} = d(j, h) - d(j, t)$$

What is the problem with PAM

- PAM is more robust than k-means in the presence of noise and outliers because a medoid is less influenced by outliers or other extreme values than a mean
- PAM works efficiently for small data sets but does not scale well for large data sets.
 - $O(k(n-k)^2)$ for each iteration

where n is # of instances, k is # of clusters

➔ Sampling based method,

CLARA (Clustering LARge Applications)

- CLARA (Kaufmann and Rousseeuw in 1990)
 - Built in statistical analysis packages, such as S+
- It **draws multiple samples** of the data set, applies PAM on each sample, and gives the best clustering as the output
- Strength: deals with larger data sets than PAM
- Weakness:
 - Efficiency depends on the sample size
 - A good clustering based on samples will not necessarily represent a good clustering of the whole data set if the sample is biased

Overview of Clustering Methods

	General Characteristics
Partitioning methods	<ul style="list-style-type: none">– Find mutually exclusive clusters of spherical shape– Distance-based– May use mean or medoid to represent cluster center– Effective for small- to medium-size data sets
Hierarchical Methods (aside)	<ul style="list-style-type: none">– Clustering is a hierarchical decomposition (i.e., multiple levels)– Cannot correct erroneous merges or splits– May incorporate other techniques like microclustering or consider object “linkages”
Density-based methods	<ul style="list-style-type: none">– Can find arbitrarily shaped clusters– Clusters are dense regions of objects in space that are separated by low-density regions– Cluster density: Each point must have a minimum number of points within its “neighbourhood”– May filter out outliers
Grid-based Methods (aside)	<ul style="list-style-type: none">– Use a multiresolution grid data structure– Fast processing time (typically independent of the number of data objects, yet dependent on grid size)

- In clustering, data are
 - similar to one another within the same cluster, and
 - dissimilar to the data in other clusters.
- Cluster analysis can be used as
 - a stand-alone tool to gain insight into the data distribution or
 - can serve as a pre-processing step for other data mining tasks.
- Partitioning method:
 - iterative relocation technique: k-means and k-medoids.
 - K-medoid is efficient in presence of noise and outliers.
- Density based method
 - Discover clusters of arbitrary shape,
 - good at handling noise,
 - requires only one scan
 - sensitive to density parameters (needed to be set manually)

- Han et al.'s book
 - Chapter 2.1, 2.4 and 10
- Readings
 - [An Introduction to k-means](#)
 - [K-means++](#)